



Some definitions

- A **processor** is a small chip that responds to and processes the basic instructions that drive a computer. The term *processor* is used interchangeably with the term **central processing unit** (CPU)
- **Core**: The smallest compute unit that can run a program
- **Socket**: A compute unit, packaged as one and usually made of a single chip often called processor. Modern sockets carry many cores (10, 14, or 20, 24, 28, etc. on most servers)
- **Node**: A stand-alone computer system that contains one or more sockets, memory, storage, etc. connected to other nodes via a fast network interconnect.


From a CPU to a Cluster



CPU
(Has Multiple cores)




Socket
(Can have Multiple CPUs)




Motherboard



Compute Node

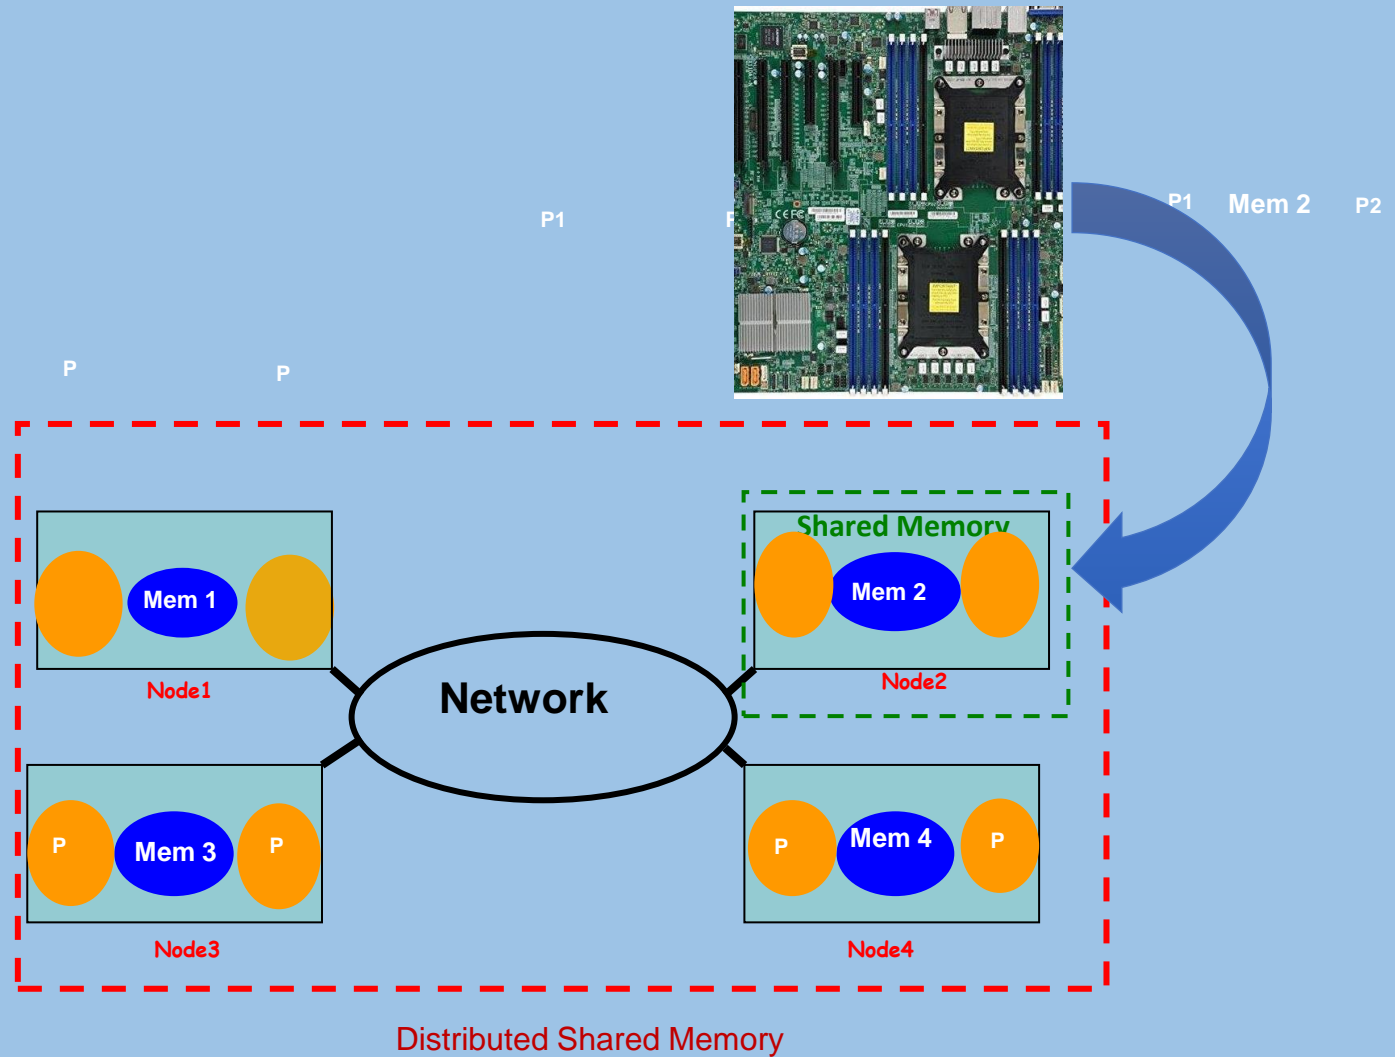


Rack



Cluster

Distributed vs Shared Memory Systems



How to Run jobs

**Either Interactively or submitting
jobs to a queue using Slurm**

Slurm: Some key terms to remember

A **job** is the resources you are using and the code you are running

The **queue** in Slurm is all RUNNING and all PENDING jobs
To see every job in the queue on Oscar, use the command

```
squeue
```

To see your jobs in the queue

```
squeue -u <cnetid>
```

```
or
```

```
myq
```

What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

This **#!** is a shebang

It tells operating system to use
/bin/bash
with this script

What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

is a comment

everything after # is ignored by
bash

What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

#SBATCH is a directive

It is a comment in Bash

#SBATCH is only relevant to slurm:
sbatch my_script.sh

What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

#SBATCH is a directive

It is a comment in Bash

#SBATCH is only relevant to slurm:
sbatch my_script.sh

To comment out directives, break
the pattern, e.g.

```
##SBATCH
```

```
# SBATCH
```

What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

Instructions for Slurm must go at the top of the script

Any #SBATCH lines you put after your program will be ignored

What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash

# Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

Slurm has some variables you can use. %j is the job number. When the job runs %j will be expanded to the job number. In this example %j is used in the output file and error file names:

MyJob-13571056.out
MyJob-13571056.err

%j is unique. By using %j in your filenames you guarantee a unique file name, which means you won't accidentally overwrite previous output.

What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash
```

```
# Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

=> Time your job is allowed to run

=> Number of nodes to run on

=> Number of cores on each node to use

=> Memory per cpu => 2000Mb or 2Gb

=> Name of the job.

=> Job output file behaves as stdout for the code.

=> Error file. behaves as stderr for the code.

=> Load any modules you need for your application

=> run the code you want

Running batch jobs using a Submission Script

- A simple job submission script (saved as python.sbatch):

Slurm {
Your Job {

```
#!/bin/bash
#SBATCH --job-name=first_python_job
#SBATCH --output=first_python_job_%j.out
#SBATCH --error=first_python_job_%j.err
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000M
#SBATCH --partition=broadwl
#SBATCH --reservation=kicpworkshop-cpu
#SBATCH --time=00:30:00
module load python
python hello_world.py
echo "job finished at `date`"
```

- To submit the above script:
 - **sbatch** python.sbatch

Exercise;
Ex-1 in Repo

How to submit OpenMP jobs?

```
#!/bin/bash

#Here is a comment
#SBATCH --time=1:00:00
#SBATCH --partition=broadwl
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=8
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
export OMP_NUM_THREADS=8
#Run your code
./my_executable
```

Specify number of cores > 1.

OMP_NUM_THREADS is an environment variable.

Exercise;
Ex-2 in Repo

How to submit Parallel MPI jobs?

```
#!/bin/bash
```

```
#Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
#SBATCH --partition=broadwl
```

```
#SBATCH --nodes=4
```

```
#SBATCH --ntasks-per-node=8
```

```
#SBATCH --mem-per-cpu=2000
```

```
module load openmpi
```

```
module load <module name>
```

```
#Run your code
```

```
mpirun ./my_executable
```

Specify number of nodes > 1.

Specify number of cores >= 1.

Load OPENMPI MPI library or IntelMPI

Exercise;
Ex-3 in Repo

How to submit GPU jobs?

```
#!/bin/bash
```

```
#Here is a comment
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --partition=gpu2
```

```
#SBATCH --gres=gpu:1
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=8
```

```
#SBATCH --mem-per-cpu=2000
```

```
#SBATCH --job-name=MyJob
```

```
#SBATCH --output= MyJob-%j.out
```

```
#SBATCH --error=MyJob-%j.err
```

```
module load <module name>
```

```
#Run your code
```

```
./my_executable
```

Specify partition `gpu2`


Specify number of gpus, like `gpu:1`

Exercise;
Ex-4 in Repo

How to submit MPI + GPU jobs?

CUDA Aware MPI

```
mpirun -np 4 ./myapp <args>
```



How to submit MPI + GPU jobs?

Compilation

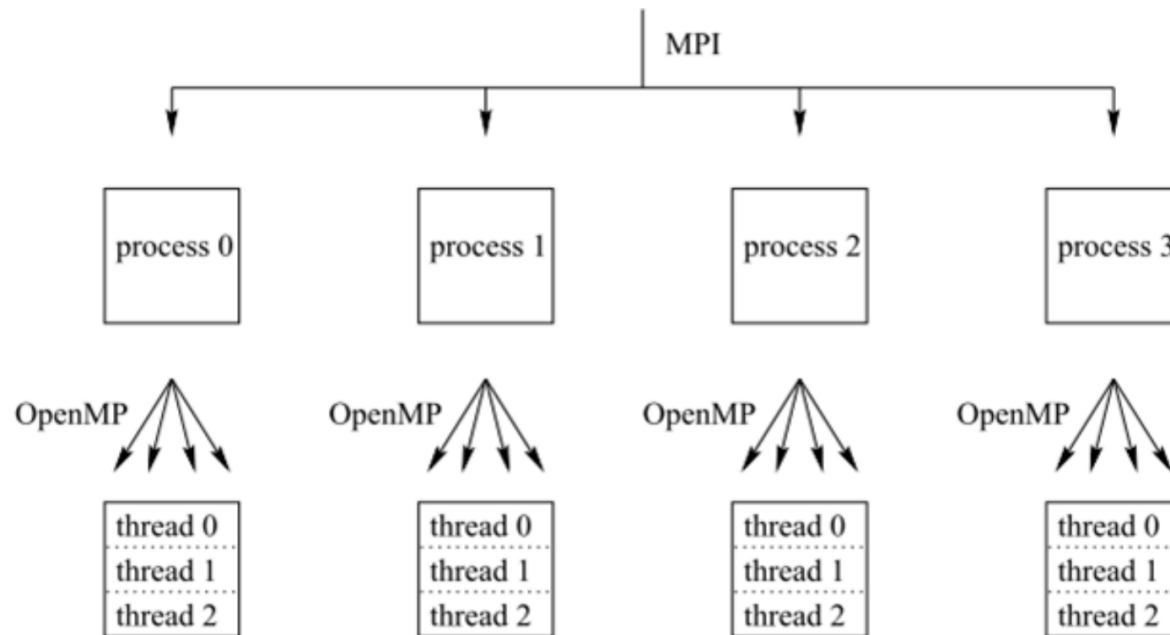
```
#!/bin/bash
module load openmpi/3.1.2
module load cuda/10.1
# Compiling the device code
nvcc -c dev.cu
# Compiling the host code
mpicc -c hostname.c
# Linking the host and device code
mpicc -o HostMap dev.o hostname.o -lcudart
# Submitting the job as batch script
sbatch mpijob.sh
```

Job Submission

```
#!/bin/bash
#SBATCH -t 00:30:00
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=4
#SBATCH --partition=gpu2
#SBATCH --gres=gpu:2
#SBATCH --job-name=MyJob
#SBATCH --output=MyJob-%j.out
#SBATCH --error=MyJob-%j.err
#SBATCH --qos=stafftest
mpirun ./HostMap
```

Exercise;
Ex-5 in Repo

How to submit MPI + OpenMPI jobs?



```
#!/bin/bash
#SBATCH --job-name=hybrid
#SBATCH --output=hybrid_%j.out
#SBATCH --error=hybrid_%j.err
#SBATCH --time=00:10:00
#SBATCH --ntasks=4
#SBATCH --cpus-per-task=8
#SBATCH --partition=broadwl
#SBATCH --constraint=edr
#SBATCH --qos=stafftest

# Load the default OpenMPI module.
module load openmpi

# Set OMP_NUM_THREADS to the number of CPUs per task we asked for.
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

mpirun ./mpomp
```

Exercise;
Ex-5a in Repo

How to submit array based jobs

```
#!/bin/bash
```

```
# Job Name
```

```
#SBATCH --job-name=arrayjob
```

```
# Walltime requested
```

```
#SBATCH --time=0:10:00
```

```
#Add partition
```

```
#SBATCH --partition=broadwl-lc
```

```
# Provide index values (TASK IDs)
```

```
#SBATCH --array=1-16
```

```
# Use '%A' for array-job ID, '%J' for job ID and '%a' for task ID
```

```
#SBATCH --error=maths%A-%a.err
```

```
#SBATCH --output=maths%A-%a.out
```

```
# single core
```

```
#SBATCH --ntasks-per-node=1
```

```
#SBATCH --mem-per-cpu=2000
```

```
# Use the $SLURM_ARRAY_TASK_ID variable to provide different inputs for each job
```

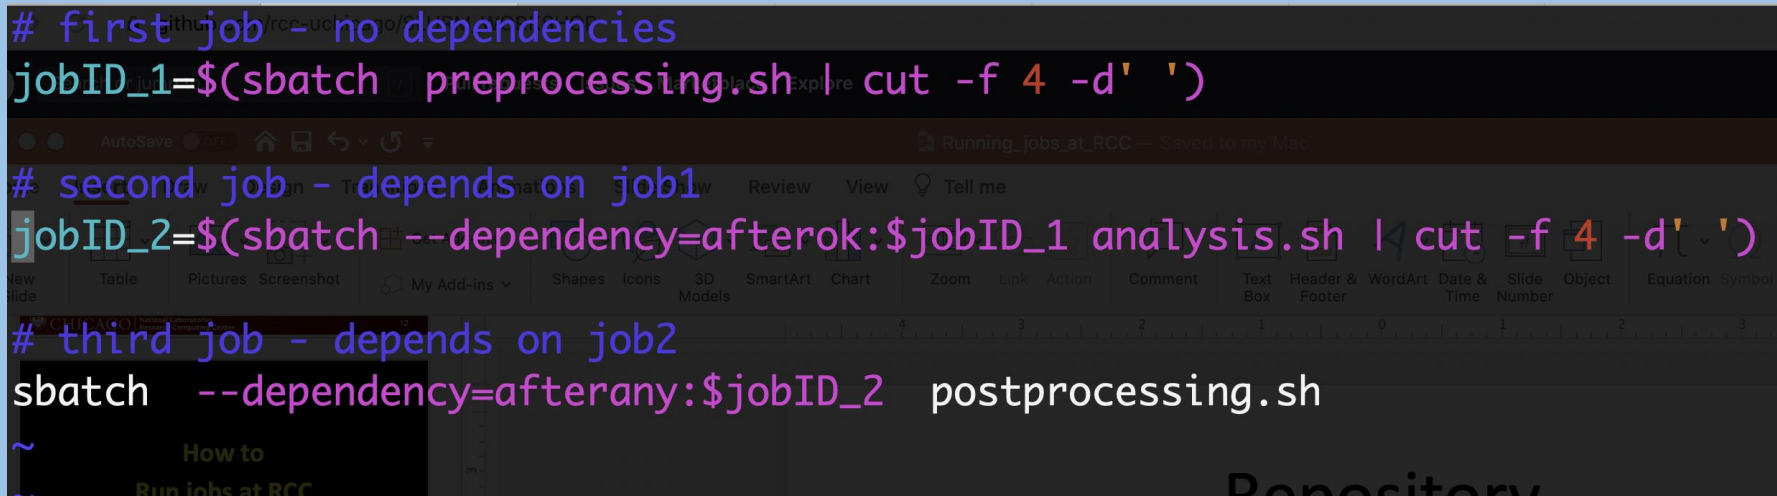
```
input=$((SLURM_ARRAY_TASK_ID*1000+2))
```

Exercise;
Ex-6 in Repo

How to submit dependent jobs

SLURM Rule:

sbatch --dependency=type:job_id jobfile



```
# first job - no dependencies
jobID_1=$(sbatch preprocessing.sh | cut -f 4 -d' ')

# second job - depends on job1
jobID_2=$(sbatch --dependency=afterok:$jobID_1 analysis.sh | cut -f 4 -d' ')

# third job - depends on job2
sbatch --dependency=afterany:$jobID_2 postprocessing.sh
```

How to submit dependent jobs

after	This job can begin execution after the specified jobs have begun execution
afterany	This job can begin execution after the specified jobs have terminated.
aftercorr	A task of this job array can begin execution after the corresponding task ID in the specified job has completed successfully
afternotok	This job can begin execution after the specified jobs have terminated in some failed state
afterok	This job can begin execution after the specified jobs have successfully executed
singleton	This job can begin execution after any previously launched jobs sharing the same job name and user have terminated

Exercise;
Ex-7 in Repo

How to submit Parallel batch jobs

```
#!/bin/sh
```

```
#SBATCH --time=01:00:00
```

```
#SBATCH --partition=broadwl
```

```
#SBATCH --ntasks=28
```

```
#SBATCH --mem-per-cpu=2G # NOTE DO NOT USE THE --mem= OPTION
```

```
# Load the default version of GNU parallel. module load parallel
```

```
# set the max number of processes (which determine the max per processor)
```

```
ulimit -u 10000
```

```
# This specifies the options used to run srun. The "-N1 -n1" options are
```

```
# used to allocates a single core to each task.
```

```
srun="srun --exclusive -N1 -n1"
```

```
#Run GNU parallel
```

```
parallel="parallel --delay 0.2 -j $SLURM_NTASKS --joblog runtask.log --resume"
```

```
# Run a script, runtask.sh, using GNU parallel and srun.
```

```
$parallel "$srun ./runtask.sh arg1:{1} > runtask.sh.{1}" ::: {1..128}
```

```
# Note that if your program does not take any input, use the -n0 option to call the  
parallel command: # # $parallel -n0
```

```
"$srun ./run_noinput_task.sh > output.{1}" ::: {1..128}
```


How to submit Parallel batch jobs

Exercise;
Ex-8 in Repo

What resources should I ask for?

This depends on the code you are running

What resources should I ask for?

This depends on the code you are running

What resources should I ask for?

This depends on the code you are running

Nodes/Cores

- Question: is your code parallel? You will need to find out if your code can
 - Run on multiple cores? Run across multiple nodes?
 - Check if your code is threaded, multiprocessor, MPI
- Question: Is your code serial?
 - This means it can only make use of one core

What resources should I ask for?

This depends on the code you are running

Wall Time

Make an estimate of your job run and add a bit.

e.g. if think your code will take an hour, give it 1 hour and 30 min

- If your job runs out of time, your job will be killed, so
- be accurate with your estimate without going below.

What resources should I ask for?

This depends on the code you are running

Memory

For memory, this can take some trial and error. You can ask for a lot, then measure your usage. If you have asked for more memory and then reduce your memory with the next job.

- To ask for all the memory available on a node, use `#SBATCH --mem=0`

What if I need an entire node or specific features?

Add this in your batch script

```
#SBATCH --exclusive
```

Add this in your batch script

```
#SBATCH --constraint=v100
```

How do I know the features of the node to use with #SBATCH -constraint?

nodestatus

nodestatus

Status of nodes:

Features

STATUS CORES IN USE

MEM IN USE PURPOSE

midway2-0002	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	mix	16	57.1%	24GB	41.9%	broadwl
midway2-0003	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	alloc	28	100%	5GB	9%	broadwl
midway2-0004	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	alloc	28	100%	15GB	26.7%	broadwl
midway2-0005	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	alloc	28	100%	7GB	12.4%	broadwl
midway2-0006	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	mix	25	89.2%	16GB	28.7%	broadwl
midway2-0007	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	mix	16	57.1%	4GB	8.2%	broadwl
midway2-0008	28-core	58GB	tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b	mix	19	67.8%	19GB	33.1%	broadwl

What resources should I ask for?

This depends on the code you are running

GPUs If your code is built to use gpus you can submit to the gpu partition. To request 1 gpu:

```
#SBATCH -p gpu2 --gres=gpu:1
```

What resources did my job actually use?

It is good practice to occasionally check what resources your job is using. For example if you are going to be submitting hundreds of similar jobs, you may save yourself a lot of waiting time in the queue by checking that you are not over requesting resources.

Why did my job fail?

- Ran out of Memory
- Ran out of Time
- Problem in your submit script
- Problem with your code
- Node failure

Why did my job fail?

- Ran out of Memory
- Ran out of Time
- Problem in your submit script
- Problem with your code
- Node failure

You can fix these

Why did my job fail?

- Ran out of Memory
- Ran out of Time
- Problem in your submit script
- Problem with your code
- Node failure


You can fix these

Job Priority

- Priority is calculated using a **Fairshare** algorithm
- Fairshare is a function of
 - Requested wall clock, memory, nodes/cores, etc.
 - Length of time in queue
 - Number of jobs in a time window and per PI group
 - Backfill
 - Etc.

Job Priority

Backfill




the x axis is time

time

Job Priority

Backfill

Why has someone else's job started before mine?




The scheduler can fit smaller jobs in before big jobs are due to start

Job Priority

Backfill


Why has someone else's job started before mine?



Job Priority

Backfill

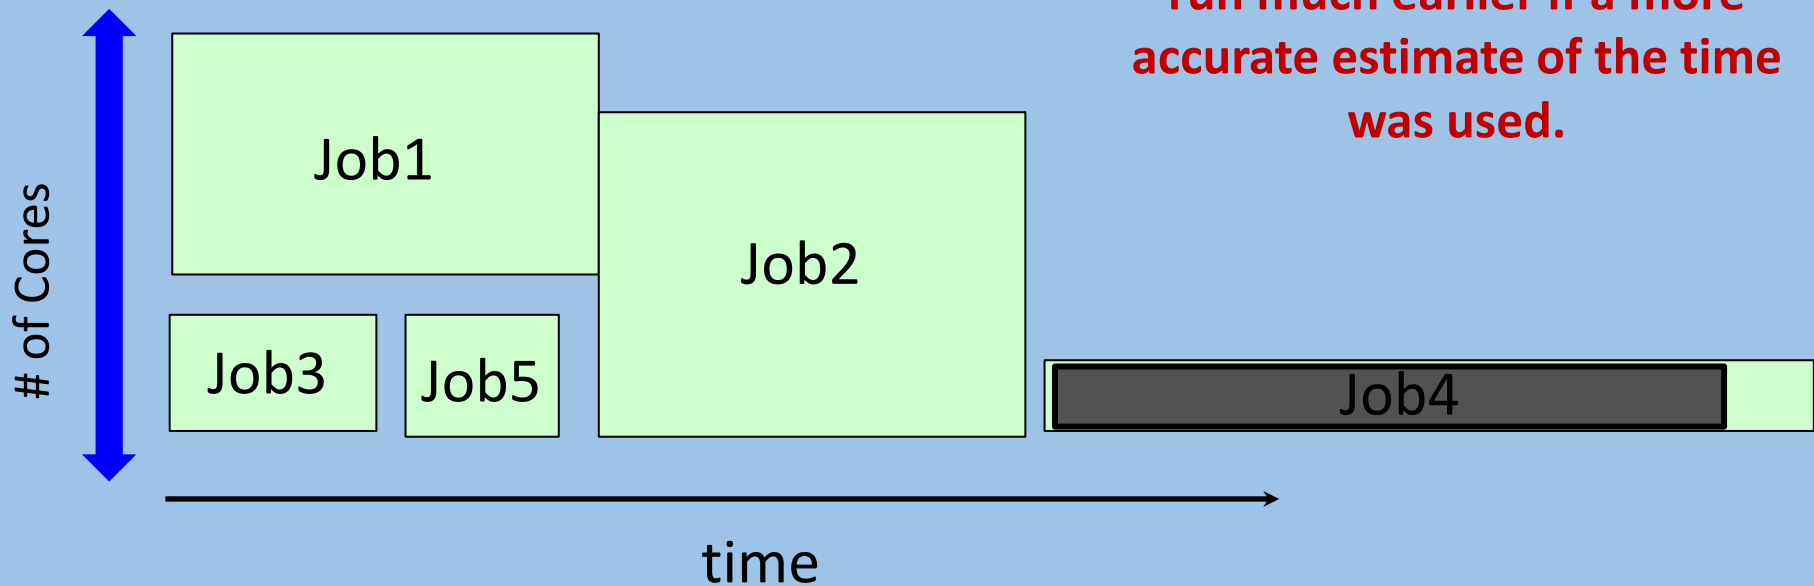
Why has someone else's job started before mine?



Job Priority

Backfill

Why has someone else's job started before mine?



Job submission and monitoring

SLURM Commands

Command	Description
<code>sbatch script.sbatch</code>	Submits <code>script.sbatch</code> job script
<code>squeue -u \$USER or myq</code>	Reports the status of your jobs
<code>sacct -u \$USER</code>	Displays accounting data for your job(s)
<code>scancel jobid</code>	Cancels a running job or removes it from the queue
<code>scontrol show job jobid or jobinfo</code>	Displays details of a running job